

A Formal Semantics for Skip’s Reactive `reduce` Combinator

Cristiano Calcagno

November 2025

Abstract

We present a formal semantics for the `reduce` combinator in Skip [1], a reactive programming framework that maintains aggregated views of collections with automatic incremental updates. Skip exposes reducers as **first-class combinators**: a reducer $R = (\iota, \oplus, \ominus)$ is a reusable, composable value—not a callback tied to a specific engine context—that developers can apply uniformly across arbitrary reactive pipelines. Unlike streaming systems (Flink, Kafka Streams) that provide add/remove callbacks bound to specific operators, and unlike FRP and incremental computation frameworks that hide inverses internally, Skip surfaces the algebraic structure as a user-facing programming construct. This paper formalizes Skip’s correctness condition—providing a user-facing specification not found in prior systems—and proves that incremental correctness holds *if and only if* the reducer is well-formed (i.e., \ominus is the inverse of \oplus), giving Skip users a precise specification for writing correct custom reducers.

1 Introduction

Skip [1] is a reactive programming framework that maintains derived views of collections with automatic incremental updates. When the underlying data changes, Skip efficiently propagates updates to all dependent computations without manual intervention. A central operation in Skip is the `reduce` combinator, which computes a summary (or *view*) for each key in a collection by folding over its associated values—for example, computing the sum, count, or minimum.

The key to efficient updates is that Skip’s `reduce` supports *retractions*: when values are removed from a collection, the reducer can incrementally update the accumulated result rather than recomputing from scratch. Skip exposes this capability through a *reducer* abstraction $R = (\iota, \oplus, \ominus)$, where ι is an initial value, \oplus is an add operation, and \ominus is a remove operation. This allows $O(1)$ updates per change, rather than $O(n)$ recomputation.

Skip’s user-facing combinators. A distinguishing feature of Skip is that reducers are exposed as **user-facing combinators**—first-class programming constructs that developers use directly to build custom reactive services. Rather than being internal implementation details hidden from users, Skip allows developers to define their own reducers for domain-specific aggregations. Skip’s documentation specifies an informal correctness condition: the result of applying the runtime’s sequence of remove/add calls must equal recomputing from scratch. This design enables extensibility (custom aggregations), composability (reducers combine with other Skip combinators), and—when the condition is satisfied—correctness guarantees.

What distinguishes Skip’s design. Three aspects distinguish Skip’s reducer from related systems:

1. **First-class combinator, not callback.** Unlike Flink’s `retract` methods or Kafka Streams’ `Subtractor` interfaces—which are callbacks tied to specific engine contexts—Skip’s reducer is a standalone, reusable value applicable to any reactive collection.
2. **User-facing inverse.** Unlike FRP libraries (which provide only append-style folds) and incremental computation frameworks (which handle inverses internally), Skip exposes \ominus as an explicit user-defined operation.
3. **Formal correctness.** While systems like DBSP and differential dataflow rest on rigorous algebraic foundations internally, no prior system provides a formal specification that users can apply when defining their own reducers; this paper fills that gap.

This paper. We formalize Skip’s reducer abstraction and its correctness condition. Our main result (Theorem 2) proves that incremental correctness holds *if and only if* the reducer is well-formed—that is, \ominus is the inverse of \oplus . This provides Skip users with a precise specification: satisfy the well-formedness condition, and your custom reducer is guaranteed to work correctly with Skip’s incremental update mechanism.

Synthesis of ideas. Our formalization builds on Skip’s design and synthesizes ideas from multiple domains, which anchor the rest of the paper:

- From **incremental databases and streaming systems** (distributive, algebraic, holistic aggregates [6]), we carry over the aggregate-class taxonomy and view maintenance perspective, and the insight that invertibility yields $O(1)$ per-change maintenance for suitable aggregates.
- From **Skip’s reactive runtime**, we take the user-facing combinator contract: reducers are first-class, composable values whose invertibility the runtime can exploit, and which otherwise fall back safely to recompute.

- From **formal methods**, we provide a complete characterization: not merely sufficient conditions for correctness, but a precise *if-and-only-if* theorem that fully characterizes when Skip’s incremental updates are correct.

Contributions. This paper formalizes Skip’s `reduce` combinator. We provide:

- A denotational semantics for `reduce` as a derived view (Section 3)
- A formal model of deltas and Skip’s incremental update procedure (Section 4)
- A precise well-formedness condition and proof that it is both necessary and sufficient for correctness (Section 5)
- A complexity contract: well-formed reducers admit $O(1)$ per-key updates, while partial reducers fall back to recomputation (Sections 4 and 6)
- Concrete examples including sum, count, and min reducers (Section 6)

We also position Skip’s design relative to other streaming and reactive systems (Section 7) and analyze the complexity benefits of incremental updates.

2 Preliminaries

Let K be a set of keys, V a set of values, and A a set of accumulator values. For a set V , we write $\mathcal{M}(V)$ for the set of finite multisets over V ; we use \uplus and \setminus for multiset union and multiset difference, respectively, and write $M \subseteq N$ for multisets when every element has multiplicity in M less than or equal to its multiplicity in N .

Definition 1 (Collection). *A collection is a function $C : K \rightarrow \mathcal{M}(V)$. We write $C(k)$ for the multiset of values associated with key k .*

Example 1 (Purchases per user). *As a running example, let $K = \{\text{user1}, \text{user2}\}$ and $V = \mathbb{Z}$. We model purchases per user and care about which purchase amounts occurred, not about the order in which they happened. A collection $C : K \rightarrow \mathcal{M}(V)$ might be given by*

$$C(\text{user1}) = \{30, 30, 50\} \quad \text{and} \quad C(\text{user2}) = \{70\}.$$

Here $C(\text{user1})$ is a multiset where 30 has multiplicity 2 and 50 has multiplicity 1, meaning that user1 made two purchases of \$30 and one purchase of \$50 (in some order). This illustrates that multiple purchases can have the same amount, and that multisets record these multiplicities while abstracting away from order.

2.1 Commutative Operations

Definition 2 (Pairwise Commutative Operation). *Let $\star : A \times V \rightarrow A$ be an update operation. We say that \star is pairwise commutative if*

$$\forall a \in A, v_1, v_2 \in V. (a \star v_1) \star v_2 = (a \star v_2) \star v_1.$$

2.2 Folds

Definition 3 (Fold over Sequence for an Operation). *Let $\star : A \times V \rightarrow A$ be an update operation and let $s = [v_1, \dots, v_n]$ be a finite sequence of elements of V . For any $a \in A$ we define:*

$$\text{fold}_\star^{\text{seq}}(a, []) = a \quad \text{and} \quad \text{fold}_\star^{\text{seq}}(a, v_1 :: s') = \text{fold}_\star^{\text{seq}}(a \star v_1, s').$$

When a distinguished initial element $\iota \in A$ is understood from context, we write $\text{fold}_\star^{\text{seq}}(s)$ for $\text{fold}_\star^{\text{seq}}(\iota, s)$.

Theorem 1 (Characterisation of Multiset Fold). *Let $\star : A \times V \rightarrow A$ be an update operation. The following are equivalent:*

1. *For all $a \in A$, $M \in \mathcal{M}(V)$ and any two finite sequences s_1, s_2 enumerating M (with multiplicity), we have*

$$\text{fold}_\star^{\text{seq}}(a, s_1) = \text{fold}_\star^{\text{seq}}(a, s_2).$$

That is, folding depends only on the multiset of elements, not on their enumeration.

2. *The operation \star is pairwise commutative in the sense of Definition 2.*

Proof. Sketch: For $(2 \Rightarrow 1)$, one shows first that swapping two adjacent elements in a sequence does not change the fold, using pairwise commutativity of \star . Since any permutation of a finite sequence can be written as a product of adjacent transpositions, it follows that the fold depends only on the underlying multiset. For $(1 \Rightarrow 2)$, instantiate (1) with the two sequences $[v_1, v_2]$ and $[v_2, v_1]$ enumerating the same multiset $\{v_1, v_2\}$, and expand the definition of the fold to obtain $(a \star v_1) \star v_2 = (a \star v_2) \star v_1$. \square

Definition 4 (Fold over Multiset for an Operation). *Let $\star : A \times V \rightarrow A$ be pairwise commutative and let $M \in \mathcal{M}(V)$ be finite. For $a \in A$ and any sequence s enumerating M (with multiplicity), we set*

$$\text{fold}_\star(a, M) := \text{fold}_\star^{\text{seq}}(a, s),$$

which is well-defined by the Characterisation of Multiset Fold. If an initial element $\iota \in A$ is fixed, we abbreviate $\text{fold}_\star(M) := \text{fold}_\star(\iota, M)$.

Example 2 (Total spent over a multiset). *Continuing the purchase example, let $A = V = \mathbb{Z}$ and define $\star(a, v) = a + v$ with initial element $\iota = 0$. For the multiset $M = \{30, 30, 50\}$ of purchases for user1 we have*

$$\text{fold}_\star(0, M) = 0 + 30 + 30 + 50 = 110,$$

regardless of the order in which the elements of M are enumerated.

Lemma 1 (Fold over Union of Multisets). *Let $\star : A \times V \rightarrow A$ be pairwise commutative and let $M, N \in \mathcal{M}(V)$ be finite multisets. Then for all $a \in A$:*

$$\text{fold}_\star(a, M \uplus N) = \text{fold}_\star(\text{fold}_\star(a, M), N).$$

Proof. Choose an enumeration of $M \uplus N$ in which all elements of M appear first, followed by all elements of N . The result then follows immediately from the definition of $\text{fold}_\star^{\text{seq}}$ and the fact that fold_\star is independent of the particular enumeration. \square

Example 3 (Decomposing total spend). *With $\star(a, v) = a + v$ and $\iota = 0$, suppose we split user1's purchases as $M = \{30\}$ and $N = \{30, 50\}$ so that $M \uplus N = \{30, 30, 50\}$. Then*

$$\text{fold}_\star(0, M \uplus N) = 0 + 30 + 30 + 50 = \text{fold}_\star(\text{fold}_\star(0, M), N),$$

showing that we can compute the total spent either directly on $M \uplus N$ or by first aggregating over M and then incrementally folding in the purchases in N .

3 The Reduce Combinator

The `reduce` combinator produces a *view* of a collection by summarizing the values for each key.

Definition 5 (Reduce Combinator). *Let $\oplus : A \times V \rightarrow A$ be a pairwise commutative operation and $\iota \in A$ an initial value. Given a collection $C : K \rightarrow \mathcal{M}(V)$, the reduce combinator produces a view:*

$$\text{reduce}_{\iota, \oplus}(C) : K \rightarrow A$$

defined as:

$$\text{reduce}_{\iota, \oplus}(C)(k) = \text{fold}_\oplus(\iota, C(k))$$

That is, for each key k , we fold the operation \oplus over all values in $C(k)$, starting from ι .

Example 4 (Per-user analytics as reduce). *For the purchase collection C above, two common choices of \oplus are:*

- the sum operation $\oplus_{\text{sum}}(a, v) = a + v$ with $\iota_{\text{sum}} = 0$, giving each user their total spend;

- the count operation $\oplus_{\text{count}}(a, v) = a + 1$ with $\iota_{\text{count}} = 0$, giving each user their number of purchases.

Then, for example,

$$\text{reduce}_{\iota_{\text{sum}}, \oplus_{\text{sum}}}(C)(\text{user1}) = 110, \quad \text{reduce}_{\iota_{\text{count}}, \oplus_{\text{count}}}(C)(\text{user1}) = 3.$$

The view $\text{reduce}_{\iota, \oplus}(C)$ is a derived collection that depends on C . When C changes, the view must be updated to remain consistent. The next section addresses how to perform these updates efficiently.

4 Incremental Updates

When a collection C changes, the view $\text{reduce}_{\iota, \oplus}(C)$ must be updated. A naïve approach would recompute the fold from scratch for each affected key, requiring $O(n)$ time where n is the size of the multiset. To achieve $O(1)$ updates, we introduce a *remove* operation \ominus that can undo the effect of \oplus .

4.1 Reducers

Definition 6 (Reducer). A reducer is a triple $R = (\iota, \oplus, \ominus)$ where $\iota \in A$ is an initial value, and

$$\oplus, \ominus : A \times V \rightarrow A$$

are update operations such that both \oplus and \ominus are pairwise commutative in the sense of Definition 2. We call \oplus the add operation and \ominus the remove operation.

For a reducer $R = (\iota, \oplus, \ominus)$, we write reduce_R for $\text{reduce}_{\iota, \oplus}$.

Definition 7 (Well-Formed Reducer). A reducer $R = (\iota, \oplus, \ominus)$ is well-formed if \ominus is the *inverse* of \oplus on reachable accumulator values, that is, for all finite multisets $M \in \mathcal{M}(V)$ and all $v \in V$:

$$(\text{fold}_{\oplus}(\iota, M) \oplus v) \ominus v = \text{fold}_{\oplus}(\iota, M)$$

Example 5 (Interpreting well-formedness for sum). For R_{sum} and any multiset M of purchase amounts and value $v \in \mathbb{Z}$, well-formedness says that if we start from the total spend $\text{fold}_{\oplus}(0, M)$, then adding a purchase of amount v and immediately removing it leaves the total unchanged:

$$(\text{fold}_{\oplus}(0, M) + v) - v = \text{fold}_{\oplus}(0, M).$$

This is exactly the intuitive requirement that \ominus undo the effect of \oplus on reachable accumulator states.

In database terminology, a well-formed reducer defines an *invertible distributive aggregate* (see Section 4.2): the fold can be computed over partitions independently (distributive), and individual values can be removed from the accumulated result (invertible).

Remark 1 (Remove-Add Commutativity). *For well-formed reducers where \oplus and \ominus arise from an abelian group action on A , the following property holds automatically:*

$$\forall a \in A, v_1, v_2 \in V. (a \ominus v_1) \oplus v_2 = (a \oplus v_2) \ominus v_1$$

This ensures that the order of interleaved adds and removes does not affect the final result. All practical reducers (sum, count, product over commutative groups) satisfy this.

4.2 Aggregate Classes

The database literature [6] classifies aggregates as *distributive*, *algebraic*, or *holistic*. In our setting, a pair (ι, \oplus) defines a *distributive aggregate* when folding over a union of multisets can be decomposed into folds over the parts.

Definition 8 (Distributive Aggregate). *Let $\oplus : A \times V \rightarrow A$ be pairwise commutative and $\iota \in A$. We say that (ι, \oplus) is a *distributive aggregate* if for all finite multisets $M, N \in \mathcal{M}(V)$:*

$$\text{fold}_{\oplus}(\iota, M \uplus N) = \text{fold}_{\oplus}(\text{fold}_{\oplus}(\iota, M), N).$$

By Lemma 1, any pair (ι, \oplus) with pairwise commutative \oplus is a distributive aggregate in this sense. Moreover, a well-formed reducer $R = (\iota, \oplus, \ominus)$ (Definition 7) is precisely an *invertible distributive aggregate*: the aggregate is distributive over partitions of the multiset, and individual contributions can be removed using \ominus .

The database literature also defines *algebraic* aggregates [6]: aggregates that can be computed by maintaining a fixed number of distributive aggregates and post-processing their results. For example, average is algebraic because it can be computed from the distributive aggregates sum and count via division. In Skip, this corresponds to using a well-formed reducer with richer accumulator state (e.g., *(sum, count)* pairs) followed by a pointwise mapper to extract the final value. We illustrate this pattern in Section 6.

Finally, *holistic* aggregates [6] cannot be computed from bounded intermediate state—they potentially require access to the entire multiset. Examples include:

- **MEDIAN**: needs the full distribution to pick the middle value(s)
- **QUANTILES/PERCENTILES**: similar to median, require global ordering information
- **RANK**: depends on the position of a value within the full sorted dataset

For holistic aggregates, any exact incremental solution must maintain auxiliary state that grows with the data (e.g., the entire multiset or an order-statistic tree) in order to answer updates and queries. Skip can of course support such analyses by using richer data structures or approximations (e.g., quantile sketches), but these fall outside the constant-space, purely algebraic reducer model we formalize in this paper.

4.3 Deltas

We model updates to collections as deltas.

Definition 9 (Delta). *A delta Δ for a collection C is a pair (Δ^+, Δ^-) where:*

- $\Delta^+ : K \rightarrow \mathcal{M}(V)$ represents added values
- $\Delta^- : K \rightarrow \mathcal{M}(V)$ represents removed values

We require that $\Delta^-(k) \subseteq C(k)$ for all k (i.e., we can only remove values that exist in the collection).

Definition 10 (Delta Application). *Given a collection C and a delta $\Delta = (\Delta^+, \Delta^-)$ for C , the updated collection $C \bullet \Delta$ is defined pointwise by:*

$$C \bullet \Delta = \lambda k. (C(k) \setminus \Delta^-(k)) \uplus \Delta^+(k)$$

Intuitively, $C \bullet \Delta$ is the collection obtained by first removing all values in Δ^- from C and then adding all values in Δ^+ .

Example 6 (Changing a user's purchases). *For key $k = \text{user1}$ in the purchase collection C , suppose we remove one \$30 purchase and add a new \$20 purchase. This yields a delta with*

$$\Delta^-(k) = \{30\}, \quad \Delta^+(k) = \{20\}$$

and updated multiset

$$C'(k) = (C(k) \setminus \{30\}) \uplus \{20\} = \{30, 50, 20\},$$

which again records all purchases for user1 but with one order-irrelevant change in their amounts.

Remark 2 (Operational Construction of Deltas). *In the Skip runtime, for each key k we compute an old multiset $\text{old}(k)$ of contributing values and a new multiset $\text{new}(k)$. The delta is then constructed as:*

$$\Delta^+(k) = \text{new}(k) \setminus \text{old}(k) \quad \text{and} \quad \Delta^-(k) = \text{old}(k) \setminus \text{new}(k).$$

Note that Δ^+ and Δ^- are disjoint by construction. If C_{old} is the collection with $C_{\text{old}}(k) = \text{old}(k)$, then $\Delta^-(k) \subseteq C_{\text{old}}(k)$, so Δ is a valid delta for C_{old} . Moreover, $C_{\text{old}} \bullet \Delta = C_{\text{new}}$ where $C_{\text{new}}(k) = \text{new}(k)$.

4.4 Incremental Update

Definition 11 (Incremental Reduce). *Given the current accumulator value a_k for key k and a delta Δ , the new accumulator is computed as:*

$$\text{update}_R(a_k, \Delta, k) = \text{fold}_{\oplus}(\text{fold}_{\ominus}(a_k, \Delta^-(k)), \Delta^+(k))$$

That is, we first apply all removals $\Delta^-(k)$ to a_k using \ominus , and then apply all additions $\Delta^+(k)$ using \oplus . This is well-defined since \oplus and \ominus are pairwise commutative. Here fold_{\ominus} is the multiset fold induced by \ominus , defined exactly as in Section 2; pairwise commutativity of \ominus guarantees it is well-defined.

Example 7 (Updating a user's total spend). Let R_{sum} be the sum reducer and consider again $k = \text{user1}$ with

$$C(k) = \{30, 30, 50\}, \quad a_k = \text{reduce}_{R_{\text{sum}}}(C)(k) = 110.$$

Applying the delta above (removing \$30, adding \$20) gives

$$\text{update}_{R_{\text{sum}}}(a_k, \Delta, k) = \text{fold}_{\oplus}(\text{fold}_{\ominus}(110, \{30\}), \{20\}) = (110 - 30) + 20 = 100,$$

which matches the total obtained by recomputing from scratch on the updated multiset $C'(k) = \{30, 50, 20\}$.

5 Correctness

We now characterize exactly when incremental updates are correct. At some moment we have an *old* collection C and, for each key k , an accumulator

$$a_k = \text{reduce}_R(C)(k)$$

that agrees with the denotational semantics. A change to the collection is described abstractly by a delta Δ , yielding the *updated* collection

$$C' = C \bullet \Delta.$$

For each key k there are then two ways to obtain the *new* accumulator value:

- *Denotational recompute*: ignore a_k and compute

$$a'_k = \text{reduce}_R(C')(k),$$

i.e. start from ι and fold \oplus over the current multiset $C'(k)$.

- *Incremental update*: update the old accumulator a_k using the delta by

$$a'_k = \text{update}_R(a_k, \Delta, k),$$

i.e. first remove $\Delta^-(k)$ using \ominus , then add $\Delta^+(k)$ using \oplus .

Definition 12 (Incremental Correctness Property). A reducer R satisfies the incremental correctness property if for all collections C , all valid deltas Δ for C , and all keys k :

$$\text{reduce}_R(C \bullet \Delta)(k) = \text{update}_R(\text{reduce}_R(C)(k), \Delta, k)$$

Example 8 (Sum as a sanity check). For the sum reducer R_{sum} , the incremental correctness property simply formalizes the familiar claim that “subtracting the amounts of removed purchases and adding the amounts of new purchases yields the same total spend as recomputing the sum from scratch” for every user and every delta.

Lemma 2 (Cancellation for a Delta). *Let $R = (\iota, \oplus, \ominus)$ be a well-formed reducer and C a collection. For any valid delta $\Delta = (\Delta^+, \Delta^-)$ for C and key k , writing $M = C(k)$ and $M_0 = M \setminus \Delta^-(k)$, we have:*

$$\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), \Delta^-(k)) = \text{fold}_\oplus(\iota, M_0).$$

Proof. Write $D = \Delta^-(k)$. We proceed by induction on the size of the multiset D . If $D = \emptyset$, then $\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), D) = \text{fold}_\oplus(\iota, M)$ by definition, and $M_0 = M$, so the claim holds.

For the inductive step, suppose $D = D' \uplus \{v\}$ for some value v and submultiset $D' \subseteq M$, and assume the statement holds for D' . Since $M = M_0 \uplus D = M_0 \uplus D' \uplus \{v\}$ and \oplus is pairwise commutative, Lemma 1 (fold over union of multisets) gives:

$$\text{fold}_\oplus(\iota, M) = \text{fold}_\oplus(\text{fold}_\oplus(\iota, M_0 \uplus D'), \{v\}) = \text{fold}_\oplus(\iota, M') \oplus v$$

for the multiset $M' = M_0 \uplus D'$. Applying well-formedness with multiset M' and value v yields

$$(\text{fold}_\oplus(\iota, M') \oplus v) \ominus v = \text{fold}_\oplus(\iota, M').$$

Since $\text{fold}_\oplus(\iota, M) = \text{fold}_\oplus(\iota, M') \oplus v$, we have:

$$\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), \{v\}) = (\text{fold}_\oplus(\iota, M') \oplus v) \ominus v = \text{fold}_\oplus(\iota, M').$$

Now $M' \setminus D' = (M_0 \uplus D') \setminus D' = M_0$, so by the induction hypothesis applied to M' and D' we obtain

$$\text{fold}_\ominus(\text{fold}_\oplus(\iota, M'), D') = \text{fold}_\oplus(\iota, M' \setminus D') = \text{fold}_\oplus(\iota, M_0).$$

Since \ominus is pairwise commutative, Lemma 1 applies to \ominus as well, giving:

$$\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), D) = \text{fold}_\ominus(\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), \{v\}), D'),$$

and using the equalities above,

$$\text{fold}_\ominus(\text{fold}_\ominus(\text{fold}_\oplus(\iota, M), \{v\}), D') = \text{fold}_\ominus(\text{fold}_\oplus(\iota, M'), D'),$$

and finally

$$\text{fold}_\ominus(\text{fold}_\oplus(\iota, M'), D') = \text{fold}_\oplus(\iota, M_0),$$

which completes the proof. \square

Example 9 (Delta cancellation for total spend). *Continuing the purchase example, let R_{sum} be the sum reducer and consider key $k = \text{user1}$ with*

$$C(k) = \{30, 30, 50\}.$$

Suppose we remove one \$30 purchase and keep the other two purchases unchanged, so

$$\Delta^-(k) = \{30\}, \quad \Delta^+(k) = \emptyset,$$

and hence

$$M = C(k), \quad D = \Delta^-(k), \quad M_0 = M \setminus D = \{30, 50\}.$$

Lemma 2 states that

$$\text{fold}_\ominus(\text{fold}_\oplus(0, M), D) = \text{fold}_\oplus(0, M_0).$$

Indeed, $\text{fold}_\oplus(0, M) = 0 + 30 + 30 + 50 = 110$, so

$$\text{fold}_\ominus(110, \{30\}) = 110 - 30 = 80,$$

while

$$\text{fold}_\oplus(0, M_0) = 0 + 30 + 50 = 80.$$

Thus incrementally removing the multiset of cancelled purchases yields exactly the total spend over the remaining purchases.

The following theorem shows that the inverse property is both necessary and sufficient for incremental correctness.

Theorem 2 (Characterization of Incremental Correctness). *Let $R = (\iota, \oplus, \ominus)$ be a reducer (with pairwise commutative \oplus and \ominus). The following are equivalent:*

1. R is well-formed in the sense of Definition 7.
2. R satisfies the incremental correctness property.

Proof. We prove both directions.

(1 \Rightarrow 2): Well-formedness implies correctness.

Assume R is well-formed. Let C be a collection, $\Delta = (\Delta^+, \Delta^-)$ a valid delta for C , and k a key. Write $M = C(k)$ for the old multiset, $M' = C'(k) = (M \setminus \Delta^-(k)) \uplus \Delta^+(k)$ for the new multiset, and $a = \text{fold}_\oplus(\iota, M)$ for the old accumulator.

We must show $\text{fold}_\oplus(\iota, M') = \text{fold}_\oplus(\text{fold}_\ominus(a, \Delta^-(k)), \Delta^+(k))$.

Since $\Delta^-(k) \subseteq M$, we can write $M = M_0 \uplus \Delta^-(k)$ for some multiset M_0 . By pairwise commutativity of \oplus :

$$a = \text{fold}_\oplus(\iota, M) = \text{fold}_\oplus(\iota, M_0 \uplus \Delta^-(k)) = \text{fold}_\oplus(\text{fold}_\oplus(\iota, M_0), \Delta^-(k))$$

Let $a_0 = \text{fold}_\oplus(\iota, M_0)$. Then $a = \text{fold}_\oplus(a_0, \Delta^-(k))$.

By Lemma 2 we have

$$\text{fold}_\ominus(a, \Delta^-(k)) = \text{fold}_\oplus(\iota, M_0) = a_0.$$

Therefore:

$$\text{fold}_\oplus(\text{fold}_\ominus(a, \Delta^-(k)), \Delta^+(k)) = \text{fold}_\oplus(a_0, \Delta^+(k))$$

Since $M' = M_0 \uplus \Delta^+(k)$:

$$\text{fold}_\oplus(\iota, M') = \text{fold}_\oplus(\text{fold}_\oplus(\iota, M_0), \Delta^+(k)) = \text{fold}_\oplus(a_0, \Delta^+(k))$$

Thus both sides are equal.

(2 \Rightarrow 1): Correctness implies well-formedness.

Assume R satisfies the incremental correctness property. We must show $(a \oplus v) \ominus v = a$ for all $a \in A$ and $v \in V$.

Fix $a \in A$ and $v \in V$. We first establish the property for a of the form $a = \text{fold}_{\oplus}(\iota, M)$ for some multiset M .

Define a collection C with a single key k where $C(k) = M \uplus \{v\}$. Then:

$$\text{reduce}_R(C)(k) = \text{fold}_{\oplus}(\iota, M \uplus \{v\}) = \text{fold}_{\oplus}(\text{fold}_{\oplus}(\iota, M), \{v\}) = a \oplus v$$

Define delta Δ with $\Delta^-(k) = \{v\}$ and $\Delta^+(k) = \emptyset$. This is valid since $v \in C(k)$. The updated collection is C' with $C'(k) = M$.

By the incremental correctness property:

$$\text{reduce}_R(C')(k) = \text{update}_R(\text{reduce}_R(C)(k), \Delta, k)$$

The left side is:

$$\text{reduce}_R(C')(k) = \text{fold}_{\oplus}(\iota, M) = a$$

The right side is:

$$\text{update}_R(a \oplus v, \Delta, k) = \text{fold}_{\oplus}(\text{fold}_{\ominus}(a \oplus v, \{v\}), \emptyset) = (a \oplus v) \ominus v$$

Therefore $(a \oplus v) \ominus v = a$.

This establishes the inverse property for all a of the form $\text{fold}_{\oplus}(\iota, M)$, i.e. for all reachable accumulator values, which is exactly the condition in Definition 7. \square

Remark 3. *In many practical reducers (for example sum over integers or product over rationals), every accumulator value is reachable as $\text{fold}_{\oplus}(\iota, M)$ for some multiset M , so the definition of well-formedness above coincides with the simpler global inverse law $(a \oplus v) \ominus v = a$ for all $a \in A$, $v \in V$.*

Remark 4 (Partial Reducers in Skip). *Theorem 2 characterizes when incremental updates are correct for well-formed reducers. In Skip’s concrete API, the remove operation is allowed to be partial: a reducer’s remove function can signal “recompute from scratch” (by returning `None` in the ReScript bindings) instead of producing an updated accumulator. Such partial reducers handle cases where \ominus cannot efficiently invert \oplus —for example, computing the minimum without maintaining auxiliary state. The runtime responds by recomputing the fold from scratch for that key. Examples of partial reducers appear in Section 6.*

6 Examples

6.1 Sum Reducer

$$R_{\text{sum}} = (0, \lambda(a, v). a + v, \lambda(a, v). a - v)$$

This reducer is well-formed: addition is commutative, and subtraction is the inverse of addition.

Worked example. Suppose for key k we have $C(k) = \{3, 5, 7\}$, so the current view is $a_k = 0 + 3 + 5 + 7 = 15$. Now suppose value 5 is removed and value 2 is added, giving delta $\Delta^-(k) = \{5\}$ and $\Delta^+(k) = \{2\}$. The incremental update computes:

$$a'_k = (15 - 5) + 2 = 12$$

This matches a full recompute: $0 + 3 + 7 + 2 = 12$.

6.2 Count Reducer

$$R_{\text{count}} = (0, \lambda(a, v). a + 1, \lambda(a, v). a - 1)$$

This reducer counts the number of values, ignoring their content. It is well-formed since $(a + 1) - 1 = a$.

Worked example. For $C(k) = \{x, y, z\}$, we have $a_k = 3$. If y is removed ($\Delta^- = \{y\}$) and w is added ($\Delta^+ = \{w\}$), then:

$$a'_k = (3 - 1) + 1 = 3$$

6.3 Average Reducer (Algebraic)

The average is a classic example of an *algebraic* aggregate [6]: it can be expressed as a post-processing of a distributive aggregate over richer state. A standard encoding uses accumulator state $A = \mathbb{R} \times \mathbb{N}$ to track sum and count.

Define:

$$R_{\text{avgState}} = ((0, 0), \lambda((s, c), v). (s + v, c + 1), \lambda((s, c), v). (s - v, c - 1))$$

with accumulator state $A = \mathbb{R} \times \mathbb{N}$ tracking (sum, count). On reachable states with $c > 0$, the corresponding average is $\text{avg} = s/c$; when $c = 0$ (empty multiset), the view can be defined as a designated “no value” (e.g., `None`) or 0 depending on the application. This reducer is well-formed: addition and subtraction on sum and increment/decrement on count satisfy the inverse law on all reachable accumulator states, so Theorem 2 applies. In Skip, the average view can be implemented by first using $\text{reduce}_{R_{\text{avgState}}}$ and then applying a pointwise mapper that divides sum by count for each key.

Note that maintaining *only* the average (without count) is insufficient: to update the average when adding a value, one needs to know how many values contributed to the current average. Thus, average is genuinely an algebraic aggregate requiring auxiliary state, unlike sum or count which can be maintained with a single accumulator value.

6.4 Min Reducer (Partial)

The min reducer demonstrates why invertibility is essential for incremental updates. With accumulator $A = \mathbb{R} \cup \{+\infty\}$, the add operation is:

$$\iota = +\infty, \quad \oplus = \lambda(a, v). \min(a, v)$$

However, there is no inverse operation \ominus that works in general. Consider $C(k) = \{3, 5\}$ with $a_k = 3$.

- If we remove 5: we need $a'_k = 3$. We could define $(3 \ominus 5) = 3$ (removing a non-minimum has no effect).
- If we remove 3: we need $a'_k = 5$. But from $a_k = 3$ alone, we cannot know that 5 was the second-smallest value!

This shows min is *not* an invertible distributive aggregate: knowing only the accumulated minimum is insufficient to update the result when the minimum itself is removed.

In Skip’s implementation, min is handled as a *partial reducer*:

- The remove function signals “cannot update incrementally” (e.g., returns `None`)
- The runtime responds by recomputing from scratch: $\text{fold}_{\min}(\iota, C'(k))$
- Alternatively, one can maintain richer state (e.g., a sorted multiset of all values), making the remove operation invertible on that richer state—but this is no longer a constant-space reducer

7 Related Work

The problem of efficiently maintaining aggregations over changing data has been studied extensively. We position Skip’s design relative to streaming systems, reactive programming, and incremental computation, and describe our formal contribution.

Skip’s design: combinators vs. callbacks. Skip’s reducer $R = (\iota, \oplus, \ominus)$ is a *first-class combinator*—a reusable, composable value that can be applied to any reactive collection via reduce_R . This differs from systems that allow users to supply add/remove *callbacks* tied to specific engine internals. Callbacks in other systems are bound to particular contexts (keyed tables, windows, SQL queries); Skip’s combinator is context-agnostic and portable across different reactive pipelines.

Streaming systems with add/remove callbacks. Several streaming systems allow users to supply both add and remove logic, but not as a first-class combinator:

Apache Flink (Table API) has user-defined aggregates with `accumulate` and `retract` methods. These are lifecycle methods on a stateful object that Flink’s query planner plugs into its retraction-message algebra—not a standalone combinator.

Apache Kafka Streams provides `Aggregator` (add) and `Subtractor` (remove) interfaces for KTable aggregations. These are callbacks passed to the `aggregate()` method, tightly bound to KTable’s update mechanism.

Esper CEP supports `enter/leave` callbacks for windowed aggregation, but only within sliding-window contexts.

In all cases, the add/remove logic is configuration for a specific operator, not a reusable combinator. Moreover, none of these systems provide formal correctness conditions—they document informally what the remove function should do, but correctness is the user’s responsibility.

Systems without user-defined remove. *Apache Beam* and *Spark Streaming* do not expose user-defined inverse operations; they handle retractions internally via recomputation or diffing. *Functional reactive programming* libraries (Fran, Yampa, Reactive Banana [8]) provide append-only fold combinators (`foldp`, `accumE`) with no built-in remove. *Incremental computation* frameworks (Adaption, Jane Street’s Incremental) support removal internally via dependency tracking, but do not expose a user-facing inverse operation.

Reactive JS and FRP streams. UI-oriented reactive frameworks (React, Vue, Svelte, Solid) and stream/FRP libraries (RxJS, Elm, Yampa, Reactive Banana) expose push-based signals or events with combinators like map, filter, merge, `combineLatest`, and `scan/foldp`. They offer declarative wiring of time-varying values and typically avoid duplicate work via topological scheduling, but state updates are expressed as opaque folds; the runtime does not know or exploit user-defined inverses. As a result, invertibility and $O(1)$ maintenance for removals are left to user code, whereas Skip surfaces add/remove as a first-class contract the runtime can optimize.

Differential Dataflow and DBSP. Differential dataflow [2] and DBSP [3] provide rigorous foundations for incremental computation using Z-sets (multi-sets with integer multiplicities) and abelian groups. These frameworks *internally* ensure that all built-in aggregations have inverses, but the user typically works at a higher level (SQL, dataflow graphs) without explicitly defining \oplus and \ominus . Skip’s design surfaces this algebraic structure as a user-facing combinator; this paper provides the formal analysis.

Database view maintenance. The database literature classifies aggregates as *distributive*, *algebraic*, or *holistic* [6]. Skip’s well-formed reducers correspond to *invertible distributive aggregates*. Tangwongsan et al. [7] note that “prior work often relies on aggregation functions to be invertible” for efficient sliding-window maintenance. Yin et al. [9] require inverse functions for incremental graph aggregation. These works focus on algorithmic techniques; our contribution is to formalize correctness for Skip’s user-facing abstraction.

Summary: Skip’s design and our contribution. Skip provides a first-class reduce combinator $R = (\iota, \oplus, \ominus)$ with user-defined add and remove operations, applicable across arbitrary reactive contexts. Skip’s documentation

specifies an informal correctness condition for user-defined reducers. Our contribution is to:

1. Formalize this correctness condition as a well-formedness property
2. Prove that correctness holds *if and only if* the property is satisfied (Theorem 2)
3. Connect Skip’s design to the theory of invertible distributive aggregates

This gives Skip users a precise specification for writing correct custom reducers, backed by a complete formal characterization.

References

- [1] Skip Team. *Skip: A Reactive Programming Framework*. <https://github.com/SkipLabs/skip>, 2024.
- [2] Frank McSherry, Derek G. Murray, Rebecca Isaacs, and Michael Isard. Differential dataflow. In *Proceedings of CIDR*, 2013.
- [3] Mihai Budiu, Tej Chajed, Frank McSherry, Leonid Ryzhyk, and Val Tannen. DBSP: Automatic incremental view maintenance. In *Proceedings of VLDB*, 16(7):1601–1614, 2023.
- [4] Erik Meijer. Your mouse is a database. *Communications of the ACM*, 55(5):66–73, 2012.
- [5] Umut A. Acar, Guy E. Blelloch, and Robert Harper. Adaptive functional programming. In *Proceedings of POPL*, pages 247–259, 2002.
- [6] Ashish Gupta and Inderpal Singh Mumick. Maintenance of materialized views: Problems, techniques, and applications. *IEEE Data Engineering Bulletin*, 18(2):3–18, 1995.
- [7] Kanat Tangwongsan, Martin Hirzel, Scott Schneider, and Kun-Lung Wu. General incremental sliding-window aggregation. *Proceedings of the VLDB Endowment*, 8(7):702–713, 2015.
- [8] Conal Elliott and Paul Hudak. Functional reactive animation. In *Proceedings of ICFP*, pages 263–273, 1997.
- [9] Shufeng Yin, Huachen Zhang, Zhengyi Yang, Wentao Han, Wenguang Chen, and Yingxia Shao. GraphBolt: Dependency-driven synchronous processing of streaming graphs. In *Proceedings of ICDE*, 2022.